

# A Nontrivial Introduction to Julia For Data Science and Scientific Computing

Chris Rackauckas

University of California, Irvine

March 27, 2016

# Quick Introduction

- Third Year Math Ph.D student from MCSB.
- Main interest: stochastic analysis in theoretical (evo-) developmental biology.
  - Stochastic analysis: stochastic/nonautonomous dynamical systems and numerical methods
  - Focus on developmental signaling and cell lineages
- High-Performance Computing and Stats / Machine Learning.
- I run a blog on mathematics and scientific computing:
  - <http://www.stochasticlifestyle.com/>

# Notes About the Workshop

- I expect that you have knowledge of some programming or scripting language.
- **Advanced topics will be colored blue**
  - I do not expect everyone to get everything
  - I think it's important to expose the “lingo”
- Julia can get very low level and “CSy”
  - I will focus on the path from scripting downwards
- The projects give help for how to get started, but you will need to use the documentation
  - The ultimate goal is that you will learn how to use the documentation / Github / mailing lists / Gitter to get things done in Julia
  - Since Julia is so young, you cannot expect to get “free code” from Stack Exchange all the time!
- Do not be afraid to ask questions
- Take your time with the projects, help each other

# Outline

## ① Getting Started in Julia

- ① Why Julia? When to choose Julia?
- ② Installing/Building Julia and setting up an IDE
- ③ The package management system / Github
- ④ Basic usage: control statements, types, functions
- ⑤ Where to get help: documentation, message boards, etc.

## ② Syntax / terminology plus some innards of Julia

- ① Differences from other common languages
- ② Unique Programming Paradigms of Julia
  - ① Linear Algebra
  - ② Data-Oriented Programming
  - ③ Macro Programming
  - ④ Levels of Parallelization
  - ⑤ De-vectorization, SIMD, threading, parallelization
  - ⑥ Named Functions
  - ⑦ Subscoping
  - ⑧ LLVM

# Additional Topics

## ① Extra Projects

- ① Mathematical Modeling / Dynamical Systems
- ② Data Visualization
- ③ Multi-Node HPC (Julia's "MPI")
- ④ Interl Investigations
- ⑤ Modules and Packages
- ⑥ Language Bindings
- ⑦ Data Saving and Serialization

## ② Additional Topics

- ① Generated Functions
- ② Optimization / Machine Learning
- ③ GPGPU / Xeon Phi computing

# Why Julia?

- The purpose of Julia is to solve the “multiple-language problem of scientific computing”
  - MATLAB -> MEX, Python -> Cython, R -> ???
  - These languages require vectorization or C
- Another way to think about Julia: It's the 2015 idea of a good programming language:
  - MATLAB (1984), Python (1991), R (1993), Java (1995), C (1973), Lisp (1958)
  - Scripting and vectorization is succinct for scientific programs
  - JIT compilation works really well (Javascript)
  - Object-oriented programming has some good ideas
  - Functional/Meta programming is very useful for some tasks
  - All the other things you learned in a CS course
- Result: Julia is an LLVM-based JIT compiled language that looks like MATLAB/Python/R, but contains all the fancy guts you learned in computer science

# When to Use Julia?

- Everyday scientific computing / data science scripting
  - The best language for scripts which will need to be optimized
- Library creation for scientific computing
  - Numerical / High-performance computing libraries
  - Exploit “Vectorization” style of scientific scripting
- To speed production of things you would have done in C
- Best optimization, machine learning, automatic differentiation, stochastic differential equations and more packages
- “Data-oriented programming” structure for readability
- Unicode symbols like  $\alpha$  to enhance readability
- Replace MATLAB codes for HPCs due to licensing issues
- Use R / Python / C / MATLAB packages in one language
- The Cons of Julia
  - It's still new and rapidly evolving
  - Because of the advanced features, it's seems like it's not as beginner friendly as other scripting languages

# The Users of Julia

- ① People who want scripts to be easily optimized to run fast
- ② People who want to use fast packages/libraries written in Julia
- ③ People that need to bind together many different languages
- ④ People who want to use HPC/multi-node parallelism/GPGPU/Xeon Phi w/o OpenMP, MPI, CUDA...
  - ① People who want to use “Big Data” (distributed/shared arrays)
- ⑤ People who want to write optimized packages/libraries but want the fast development time of a scripting language
  - ① Making “libraries for scripters” that overwrite base function for automatic speedups (i.e. VML.jl, Devec.jl, ParallelAccelerator.jl)
- ⑥ People who want to mix procedural and functional programming styles
  - ① Writing compilers in Julia (i.e. ParallelAccelerator.jl)
- ⑦ Everybody!



# Example of a Julia Success: ParallelAccelerator.jl

- ParallelAccelerator.jl started in June 2015
- It is made and maintained by 6 people from IntelLabs
  - It is a compiler written in Julia
  - It takes standard scripting code (“vectorized”) and “does the optimizations an expert C++ coder would do”
    - Eliminates array bounds checking
    - Vectorizes to use SIMD / threading
    - Controls the caching
    - Many other things I may not understand
- To use it, just add “@acc” to the front of a function
  - Users don’t have to know why it works to get a MASSIVE (20x+) speedup
  - Can automatically compile code for Xeon Phi. May target OpenACC in the future for GPGPU

Summary: Written in Julia to speed up Julia code. Users don’t have to change their code to use it.

# Building / Installing Julia

- The easiest way to get started is to use the Current Release from [julialang.org](http://julialang.org)
- Note: for heavy language users you “may” want to get the nightly builds
  - As of March 2016, the nightly build is V0.5 which includes threading
- The other option is to build Julia:
  - GCC version 4.7 or later is required to build Julia
  - The default compiler on CentOS / RHEL (gcc 4.4) is too old to build Julia
  - For large systems (HPC nodes) you will want to increase the number of BLAS threads
  - Instructions for linking with MKL are found on the Julia Github

# Integrated Development Environments

- As of February 2016, the recommended IDE is Atom
  - The previous IDE, Juno via Light Table was dropped since Light Table is no longer maintained
  - The Juno team is now actively developing Atom.jl, Blink.jl, ink, atom-julia-client
  - Julia with Atom is similar to Rstudio / MATLAB
- Other popular IDEs:
  - IJulia (browser)
  - Sublime Text with IJulia
  - vim, emacs, etc.

# Example Julia Code

```

1  matrix = [1 2 3 4
2           3 4 5 6
3           4 4 4 6
4           3 3 3 3] #Define the matrix
5
6  f(x,y) = 2x+y #Create an inline function
7  @time for i=1:4,j=1:4 #Two loops in one command
8      #update matrix[i,j] = matrix[i,j] + f(i,j)
9      matrix[i,j] += f(i,j)
10 end #end statement is required for control statements like for
11
12 Pkg.add("ASCIIPLOTS") #Only use the first time!
13 using ASCIIPLOTS #Load the package
14 scatterplot(1:4,matrix[end,:],sym='*') #Plot the last row
15 scatterplot(1:4,matrix[2,:],sym='^') #Plot the second row
16 package = "ASCIIPLOTS" #Define a string
17 println("$package is a super cool package") #Print a string

```

# Example Output

```

1 0.002328 seconds (82 allocations: 5.281 KB) #Timing result
2
3 julia> scatterplot(1:4,matrix[2,:],sym='^') #Plot the second row
4
5
6      |-----^-----| 14.00
7      |
8      |
9      |
10     |
11     |
12     |
13     |
14     |
15     |
16     |
17     |
18     |
19     |
20     |
21     |
22     |
23     |
24     |
25     |^-----| 8.00
26     |-----|
27     1.00                                4.00
28
29 julia> println("$package is a super cool package") #Print a string
30 ASCIIPlots is a super cool package

```

# Basic Types of Julia

- A Vector is a one-dimensional array
- An `Array{Type,2}(n,m)` is an  $n \times m$  array of types
- `1 : n` forms a “Range type” which can be thought of as the vector  $[1, 2, \dots, n]$ .
- A dictionary `Dict` is a key-value structure
- A tuple is an immutable structure
  - Used for multiple return values
  - Very computationally efficient!
- A function is also a type

# Example Usage of Julia Types

```

1  ran = 1:5 #Define a range
2  vec = collect(ran) #Turn the range into a vector
3  vec2=Vector{Float64}(2) #Define a size 2 vector of 64-bit floats
4  vec2[1]=2; vec2[2]=3    # vec2 = [2,3], ; allows multiple commands
5  push!(vec2,5) # vec2 = [2,5]
6  mat = Array{Int64,2}(5,5) # Create an empty 5x5 matrix
7  mat2 = ones(5,5) #Create a 5x5 matrix of ones
8  Q,R = qr(mat2) #QR decompose mat2 (returns a tuple) into Q and R
9  tup = ran,vec,vec2 #Define a tuple of many different objects
10 tup[2] #This returns vec
11 dict = Dict{"a" => 1, "b" => 2, "c" => 3} #Define a dictionary
12 dict["a"] #Returns 1

```

# Defining Functions

```
1 function testFunction(x,y;z=0) #z is an optional argument
2     if z==0
3         return x+y,x*y #Return a tuple
4     else
5         return x*y*z,x+y+z #Return a different tuple
6         #whitespace is optional
7     end #End if statement
8 end #End function definition
9 a,b = testFunction(2,2) #Returns 4,4
10 a2,b2 = testFunction(2,3,z=3) #Returns 18,8
```



# The Package Management System

- Julia's package management system is Github
  - The packages are hosted as Github repositories
  - Julia packages are normally referred to with the ending ".jl"
  - Repositories register to become part of the central package management by sending a pull request to METADATA.jl
- The packages can be found / investigated at Github.com
  - Julia's error messages are hyperlinks to the page in Github

# A Very Quick Introduction to Git/Github

- Git is a common Version Control System (VCS)
- A project is a **repository (repos)**
- After one makes changes to a project, **commit** the changes
- Changes are **pulled** to the main repository hosted online
- To download the code, you **clone** the repository
- Instead of editing the main repository, one edits a **branch**
  - To get the changes of the main branch in yours, you **fetch**
  - One asks the owner of the repository to add their changes via a **pull request**
- Stable versions are cut to **releases**
- The major online server for git repositories is Github
- Github is a free service
- Anyone can get a Github account
- The code is hosted online, free for everyone to view
- Users can open **Issues** to ask for features and give bug reports to developers

# Using Julia's Package Management System

- Standard Usage

- `Pkg.add("Package")` clones the repo of the package to your system
- `"using Package"` imports the package for use in your code
- `Pkg.update()` fetches releases from METADATA.jl

- Advanced Usage

- `Pkg.build("Package")` re-compiles a package (if compilation failed)
- `Pkg.checkout("Package",branch)` checks out the branch (default: main)
- `Pkg.pin("Package")` keeps the package at a specific version / commit
- `Pkg.test("Package")` runs the Travis CL unit tests
- The tools for generating, tagging, and publishing a package exist within Julia

# Julia's Documentation

- Julia has a standard inline documentation system which is supported by Read the Docs
- Functions (not methods) are documented in this API by their input/output
- Optional function arguments are denoted by (...)

# Project 1: Setup and Using Packages/Documentation

- Get an IDE up and running
- Install the package Winston
- Go to the Github page and find the documentation
- Plot the example Stem plot
- Find the code where Stems is defined
- Search the issues for “Error when using Gadfly”
- Change to JeffBezanson’s Working branch
- Now look at the other plotting package Gadfly:
  - Find the Travis CI test coverage percentage
  - Take a look at the Gitter chat
- Go to Google Groups-Julia-users, search for posts about Gadfly
  - What are the names of the people who seem to answer all of the questions?

# Main Differences from Other Languages

<http://docs.julialang.org/en/release-0.4/manual/noteworthy-differences/>

- Julia uses 1-based indexing
- A vector of 1 slot is not a scalar
  - `Vector{Type}(size)`, `1:N`, `Float64`, etc. are **Types**
  - User types and functions are **first-class**
- A variable can hold a function (named functions)
- 2-dimensional arrays are treated as matrices
  - `*` is matrix multiplication, `A\b` solves  $Ax=b$ 
    - These operations are implemented in BLAS, Linpack
  - `.*`, `.+`, `.>`, etc. perform element-wise operations
- Arrays are dereferenced and constructed with `[]` (i.e. `A[i]`)

# Programming Paradigms of Julia 1: Linear Algebra

- Linear algebra is central to scientific computing and Julia
- Julia's linear algebra syntax follows MATLAB's closely.
  - For lots of linear algebra you can use MATLAB's documentation
- Julia's linear algebra documentation is split between two pages:
  - Multi-dimensional arrays: <http://docs.julialang.org/en/release-0.4/manual/arrays/>
  - Linear algebra: <http://docs.julialang.org/en/release-0.4/manual/linear-algebra/>

## Project 2: Linear Algebra

Use the following documentation page:

- Take two  $N \times N$  random matrices,  $A$  and  $B$ .
- Compute  $C = AB$  and  $D$ , the Hadamard product of  $A$  and  $B$  ( $A.*B$ )
- Compute the maximum and 2-norm of the last column of  $C - D$ .
- Make the  $2N \times N$  matrix which is  $C$  concatenated on top of  $D$ .
- Find the SVD-decomposition of  $D$  and plot the singular-value matrix
- Set `blas_num_threads` to different values and check the performance (parallelization) of  $A*B$ .
- Create a 100x100 sparse diagonal matrix with  $-2, 1, -2, 1, \dots$  on the diagonal



# Programming Paradigms of Julia 2: “Data-Oriented” Programming

- Julia does not implement a “full object-oriented” paradigm
  - There is no inheritance, encapsulation (except with Modules), and “data-owned methods”
  - **This is intentional and it will stay this way**
- However, types and multiple dispatch allow for a “Data-Oriented” structure
  - You can think of types as being objects
  - **Functions** call different **methods** depending on the type given
    - This is called multiple dispatch
    - Functions are a type whose fields are methods
  - User types are first class
    - Types can have a field which is another type, allowing for a form of inheritance
  - Standard types include Float64, Vector{Type}, Array{dim,Type}, **Any**, Dataframe ...
    - Functions which take in Any will compile at runtime for the type that it is given

## Project 3: “Data-Oriented” Programming

- Construct a function **solve** which takes in a matrix  $A$  and  $x_0$  and iterates  $x_{n+1} = Ax_n + \epsilon_n$ ,  $\epsilon \sim N(0, 1)$ .
- Construct a type called **model** which holds a matrix  $A$  and a vector  $x_0$ .
- Create a new method for the solve function which takes in a model and solves it 100 times and returns the mean and variance
- Create a new type called **ParameterSweep** which holds a vector of models
- Create a new method for the solve function which takes in a parameter sweep and returns a vector of tuples with  $A$ ,  $x_0$ , mean, and variance

# Programming Paradigms of Julia 3: “Macro” Programming

- Macros are implemented by the parser pre-compilation
  - Macros are methods for metaprogramming: programming the program
- Macros can be used to “expand” simple code into more complex code, while maintaining the simple code for readability
- Common macros: @time, @elapsed, @parallel, @devec, @simd, @nobounds, @thread
- Macros can block via: @macro begin #codecodecode end
- Users can write their own macros to avoid having to do repetitive programming tasks.

## Project 4: Using Macros

- Write functions which take  $N$ , generates uniformly distributed  $A$ ,  $B$ ,  $C$ , and  $D$  and solves  $A.*B.*C.*D$  in three ways:
  - $A.*B.*C.*D$
  - A for loop
  - A for loop with `@simd`, `@nubounds`
  - $A.*B.*C.*D$  with `@devec`
- Time the functions using `tic()` and `toc()`
- Time the functions with `@time`
- Use `@elapsed` to time the functions and save the values to a variable
  - Some of your first runs will be weird (the longest!). How can you fix this? Can you find out why this happened?
- Use `@benchmark` / `@benchmarks` (Extra: Try [ParallelAccelerator.jl](#))
- Find out what the native Julia implementation is for `.*`
  - Can you explain why devectorization is fastest?

# Levels of Parallelization

- SIMD parallelization is “processor-level parallelization”
  - The package is developed by Intel
  - It uses the processor’s AVX (Advanced Vector Extensions)
    - The compute cores can process more than 1 number at a time!
  - Add @simd to the beginning of an inner loop
    - The values must be able to be re-ordered
- Thread parallelization is shared-memory parallelization
  - Threads share the memory on the same computer
  - Use @thread at the beginning of a threadable loop
- Task parallelization is a general form of parallelization
  - Much easier than MPI, though the same capabilities
    - Can be used on multiple compute nodes on an HPC
  - Memory is not shared, items have to be distributed and passed
  - Either open julia with -p numberOfProcs or use addprocs(n)
    - One process is the control process, the others are workers
  - Use @parallel, pmap, etc. to distribute work to the workers

# Programming Paradigms of Julia 3: Named Functions

- User functions are first-class in Julia
  - This means that they can be modified like base functions
- Julia functions can be named, allowing inline function definitions
  - $\nabla f(x) = 5x$  is a valid function definition
- This can be useful for defining “subfunctions” which perform a repetitive task.
- Base functions can be overloaded with new functionality
  - Packages can use this to make “fast” versions
  - Example: `VML.jl`

## Project 5: Named Functions

- Extend the code from project 4 using `Vector{Float64}(0)` and `push!` to get the timings of  $A * B * C * D$  for  $N = 2^n$  with  $n = 1 : 10$ .
- Write a line of code which takes the array of times to produce a plot via Gadfly.
- Extend this to an inline function `plotTimes(timeArr,titStr)` which takes in a time array and a `titStr` and produces the appropriate plot.
- Use this function to plot the timings.
- Extend the plotting function to take in 4 arrays and plot them on the same graph using “`layer()`”

# Programming Paradigms of Julia 4: Subscoping

- Functions take the default values of their current scope.
- Variables defined in a higher scope “give their value” to lower scope by default
- **The highest scope is automatically global**
  - This is the single largest cause of slowdowns for first-time Julia users

```
1 x=5; y=7; #Defined globally
2 function scopeTest(z)
3     x += z #Changes global value
4     y = Vector{Float64}(1) #Declares a variable, local scope
5     y[1] = 2
6     return x + y + z
7 end
```



# Project 6: Scoping Caution

```
1 function f1()
2     @parallel for i = 1:100
3         var = 10
4         if var < 100
5             var = var + 1
6         end
7     end
8     var = 100 + 10
9 end
10 f1()
11 function f2()
12     @parallel for i = 1:100
13         var = 10
14         if var < 100
15             var = var + 1
16         end
17     end
18 end
19 f2()
```

# Quick Introduction to LLVM

- LLVM is the compiler the Julia on top of
  - It is an open source project that many other languages use
- Many functionalities of Julia are developed via the LLVM
  - Native LLVM compilation of CUDA kernels and Xeon Phi code will soon be part of Julia base
- You can check your LLVM code via `@code_llvm`, `@code_native`
  - One thing to look for is automatic SIMD constructs
  - Also look out for “too many extra things”

# Modeling Project: Logistic Curve (Medium)

The logistic difference equation is defined by the recursion

$$b_{n+1} = r * b_n(1 - b_n)$$

where  $b_n$  is the number of bunnies at time  $n$ . Starting with  $b_0 = .25$ , by around 400 iterations this will reach a steady state.

This steady state (or steady periodic state) is dependent on  $r$ .

Write a function which solves for the steady state(s) for each given  $r$ , and plot every state in the steady attractor for each  $r$  ( $x$ -axis is  $r$ ,  $y$  = value seen in the attractor) using PyPlot. Take  $r \in (2.9, 4)$   
Optimize this function.

# Modeling Solution

```
1 function logisticPlot()
2     tic()
3     r = 2.9:.00005:4; numAttract = 100;
4     steady = ones(length(r),1)*.25;
5     for i=1:400 ## Get to steady state
6         @devec steady = r.*steady.*(1-steady);
7     end
8     x = zeros(length(steady),numAttract);
9     x[:,1] = steady;
10    @simd for i=2:numAttract ## Grab values at the attractor
11        @inbounds @fastmath x[:,i] = r.*x[:,i-1].*(1-x[:,i-1]);
12    end
13    toc()
14    fig = figure(figsize=(20,10));
15    plot(collect(r),x,"b.",markersize=.06)
16    savefig("plot.png",dpi=300);
17 end
18 using PyPlot
19 @time logisticPlot()
```

# Data Visualization and DataFrames Project (Easy)

Follow along with the Gadfly + RDatasets demo at:

<https://github.com/timholly/gadfly/blob/master/demo.md>

# HPC Multi-Node Parallelization (Easy)

- Use the job script to call a test.jl script on the cluster:

```
1 #!/bin/bash
2 ## -N jbttest
3 ## -q <CHOOSE A QUEUE>
4 ## -pe mpich 128
5 ## -cwd                # Run in current directory
6 module load julia/0.4.3
7 julia --machinefile jbttest-pe_hostfile_mpich.$JOB_ID test.jl
```

- Make sure that all of the compute nodes are being used

```
1 hosts = @parallel for i=1:120
2     run('hostname') end
```

- Write a simple parallel loop (rand > .5 with (+) reduction), SSH into the nodes and use htop to check usage
- Information for other HPCs can be found here:  
<http://www.stochasticlifestyle.com/multi-node-parallelism-in-julia-on-an-hpc/>

# Interl Investigations Project (Quick, but Hard)

- Learn about the `@code_native` and `@which` macros here:

https:

[//sinews.siam.org/DetailsPage/tabid/607/ArticleID/744/  
Julia-A-Fast-Language-for-Numerical-Computing.aspx](https://sinews.siam.org/DetailsPage/tabid/607/ArticleID/744/Julia-A-Fast-Language-for-Numerical-Computing.aspx)

- Test the `@code_llvm` macro on some of your codes. What changes when you use `@simd`? Look for `<Vector>s`
- Try `apropos`, `?`, `@which` to find functions and check dispatches
- On one of the functions previously written, use the `fieldnames()` function to find the function pointer

# Modules and Packages (Medium)

- Wrap the functions from Project 4/5 into a module
- Export the functions in the module
- Import the module and test the functionality
- Generate a package via `Pkg.generate("Name")`
- Put the module into the package, test the package
- Setup the package to pre-compile during the first build



# Language Bindings (Hard)

## 1 C

- 1 Follow along here: <http://www.stochasticlifestyle.com/using-julias-c-interface-utilize-c-libraries/>
  - 1 If GSL is not available, Write a Hello World script in C which takes in an array of numbers and also prints the numbers.
  - 2 Write a Julia function which sorts the numbers, compile the function to a C function, and use the C calling interface to call the function from a C code

# Language Bindings Continued (Medium)

## ① Python

- ① Use the following package to import Python packages:

<https://github.com/stevengj/PyCall.jl>

## ② MATLAB

- ① Check out the following:

<https://github.com/JuliaLang/MATLAB.jl>

- ① Note that this requires you to NOT be on Windows
- ② Test some of the examples on the Github demo.
- ③ For a more advanced introduction, try porting Finite Element Method code as described here:  
<http://www.stochasticlifestyle.com/julia-ifem-1-mesh-generation-and-julias-type-system/>

## ③ R

- ① To call R with Julia, test the following package:

<https://github.com/lgautier/Rif.jl>

- ② To call Julia with R, test the following package:

<https://github.com/armgong/RJulia>

# Data Saving and Serialization (Medium)

- Try the HDF5 package for .jld and .mat files:  
<https://github.com/JuliaLang/HDF5.jl>
- Try the `serialize()` and `deserialize()` functions for high-performance serialization (note: the official mantra for future compatability is "people will try not to break the serialization format, but you shouldn't depend upon on it." Use .jld and .mat for long-term data storage)

# Generated Functions (Hard)

- Generated functions are Macros which have type information.
  - More general than macros
- Information can be found on the metaprogramming page:
  - <http://docs.julialang.org/en/release-0.4/manual/metaprogramming/>

# Optimization / Machine Learning (Medium)

- An extensive modeling framework for optimization is available via JuMP
  - This package lets one easily write problems and switch out solvers to find the best solver method
  - There are currently 13 available solvers
  - <https://jump.readthedocs.org/en/latest/>
- Other interfaces are available via Ipopt, NLOpt, etc.
- The homepage for the JuliaOpt group can be found here:  
<http://www.juliaopt.org/>
  - JuliaOpt has its own mailing list, Google Groups / Julia-opt

# GPGPU / Xeon Phi Computing (Medium)

- For tutorials on using the CUDA Runtime for GPGPU computing, see the following:
- <http://www.stochasticlifestyle.com/julia-on-the-hpc-with-gpus/>
- <http://www.stochasticlifestyle.com/multiple-gpu-on-the-hpc-with-julia/>
- The group JuliaGPU has a repository of the available packages:
- <https://github.com/JuliaGPU>
- Users can be found here:  
<https://gitter.im/JuliaGPU/meta>
- Xeon Phi linking can be found here:  
<http://www.stochasticlifestyle.com/interfacing-xeon-phi-via-julia/>
  - Native Xeon Phi linking can be found in the ParallelAccelerator.jl package (extremely new)